

GLSL

Dr. Hamish carr & Dr. Rafael Kuffner dos Anjos

Agenda



What is a shader



Shader stages



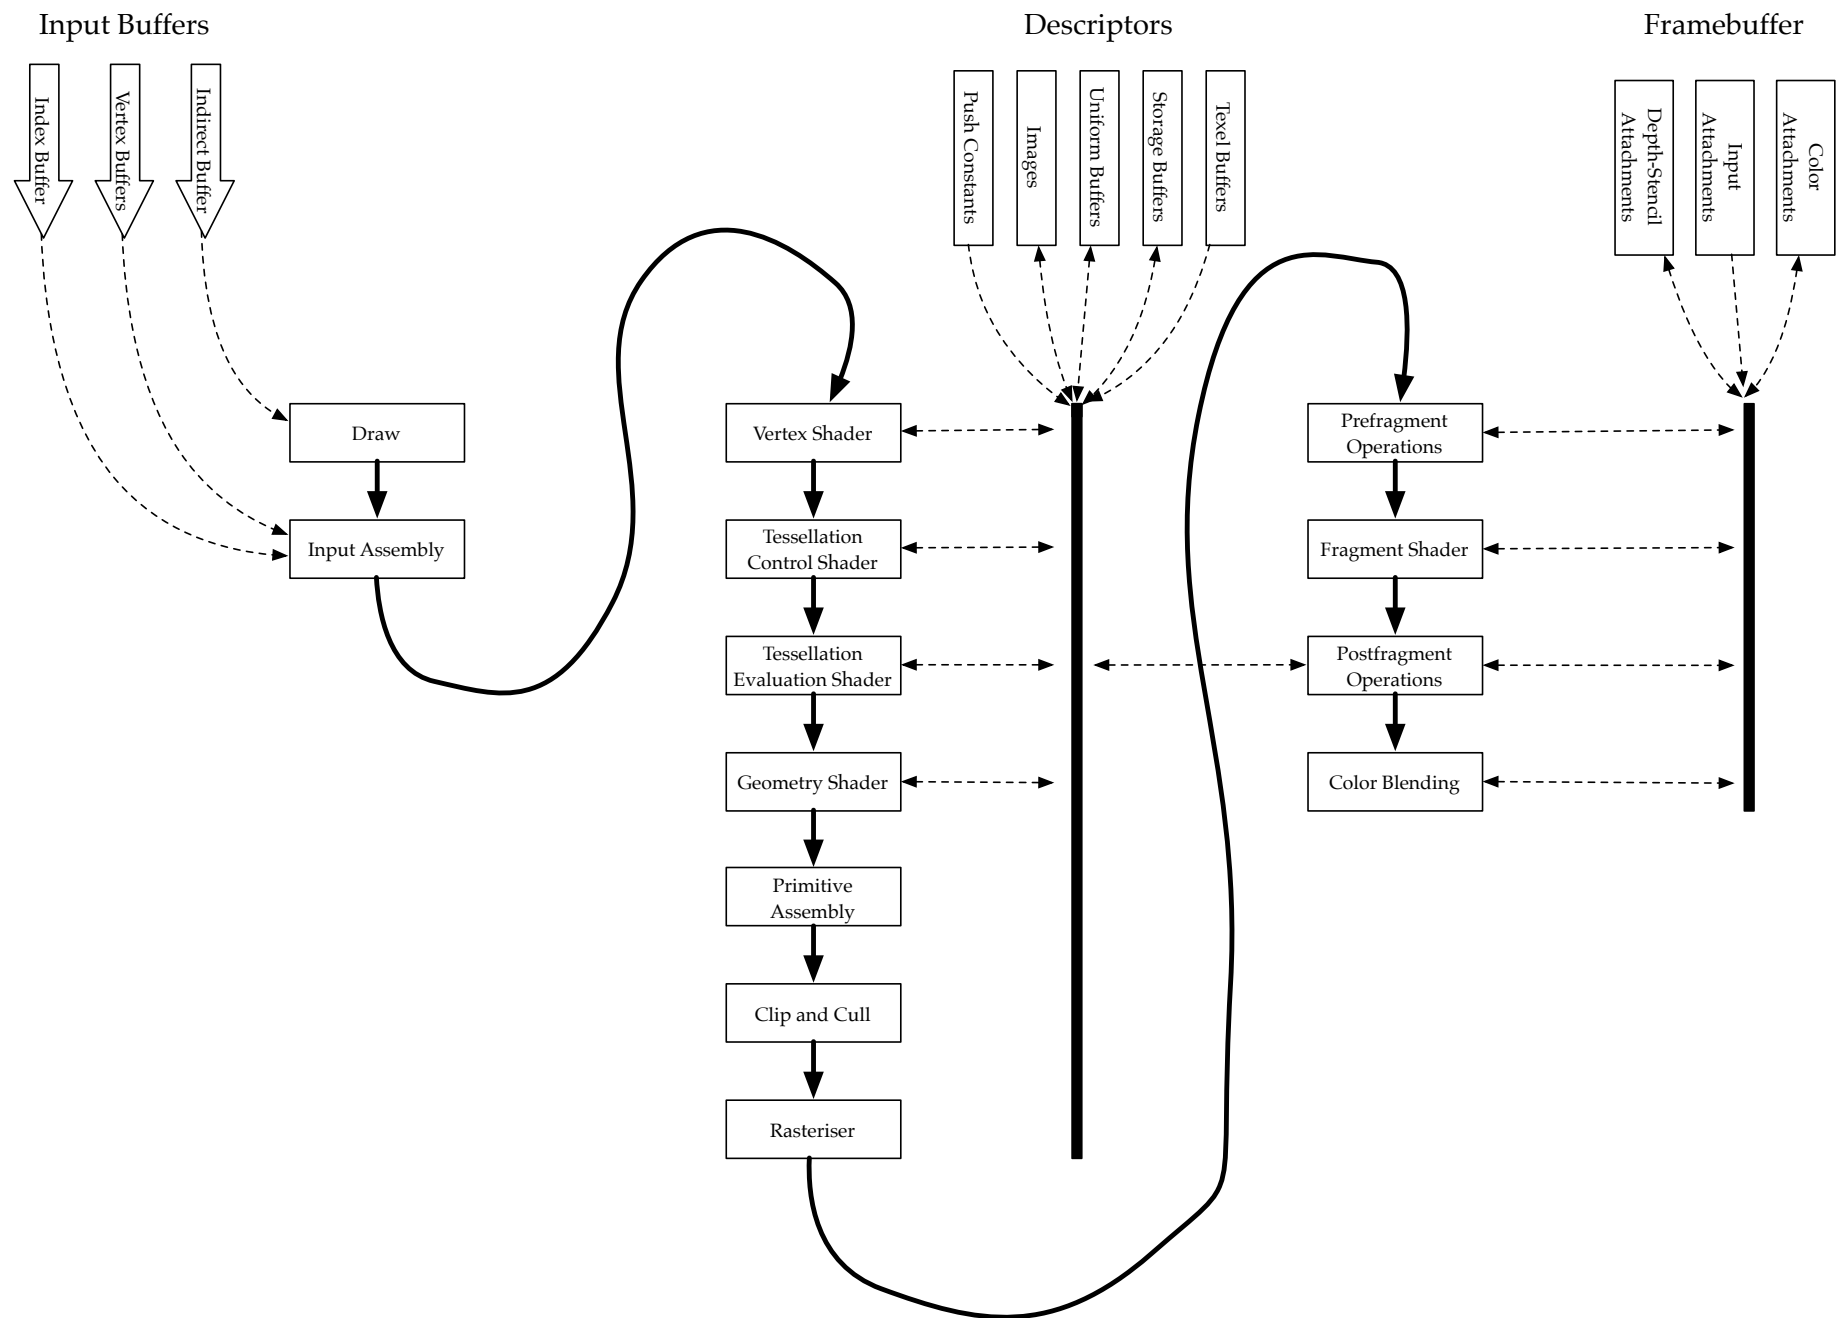
Writing a shader

What is a Shader?

- Function called per each element that needs to be processed by the pipeline
 - For each vertex -> Invoke vertex shader
 - For each fragment -> Invoke fragment shader
 - etc
- Essentially, the innards of a loop
- Applied in parallel to many data



The Vulkan Pipeline





Entrance


Early Stages

- Draw:
 - Commands enter the pipeline
- Input Assembly:
 - Reads index/vertex buffers
- Vertex Shader:
 - Transforms & processes the vertices



Tessellation Stages

- Tessellation Control Shader:
 - Generates patch tessellation commands
- Tessellation Primitive Generation:
 - Breaks patches into smaller patches
- Tessellation Evaluation Shader:
 - Sets attributes for new vertices
 - Similar to vertex shader





Geometric Stages

- Geometry Shader
 - Operates on full primitives
 - Can change primitive type
- Primitive Assembly
 - Preps vertices for rasterisation
- Clip & Cull
 - Early discard for offscreen primitives



Rasterisation

- Many options, but basically fixed function
- Rasterises & generates fragments
- Computes barycentric coordinates
- Uses them to interpolate attributes



Fragment Stages

- Prefragment Operations
 - Early discard before shading (depth, stencil)
- Fragment Assembly
 - Groups data for fragment shader
- Fragment shader
 - Code for doing shading / rendering



Fragment Processing

- Stages:
 - Scissor test – use a rectangle to clip rendering
 - Depth test – use the z-buffer to discard
 - Stencil test – use a bitmap to clip rendering
 - These can be performed *early* or *late (default)*
 - Relative to the fragment shader

Final Stages

- Postfragment Operations
 - Deferred prefragment operations
 - If Fragment shader changes data
- Color Blending
 - Updates the Framebuffer
 - Performs image processing



Writing Shaders



A shader is therefore just a small program with a `main()` routine, known as an entry point

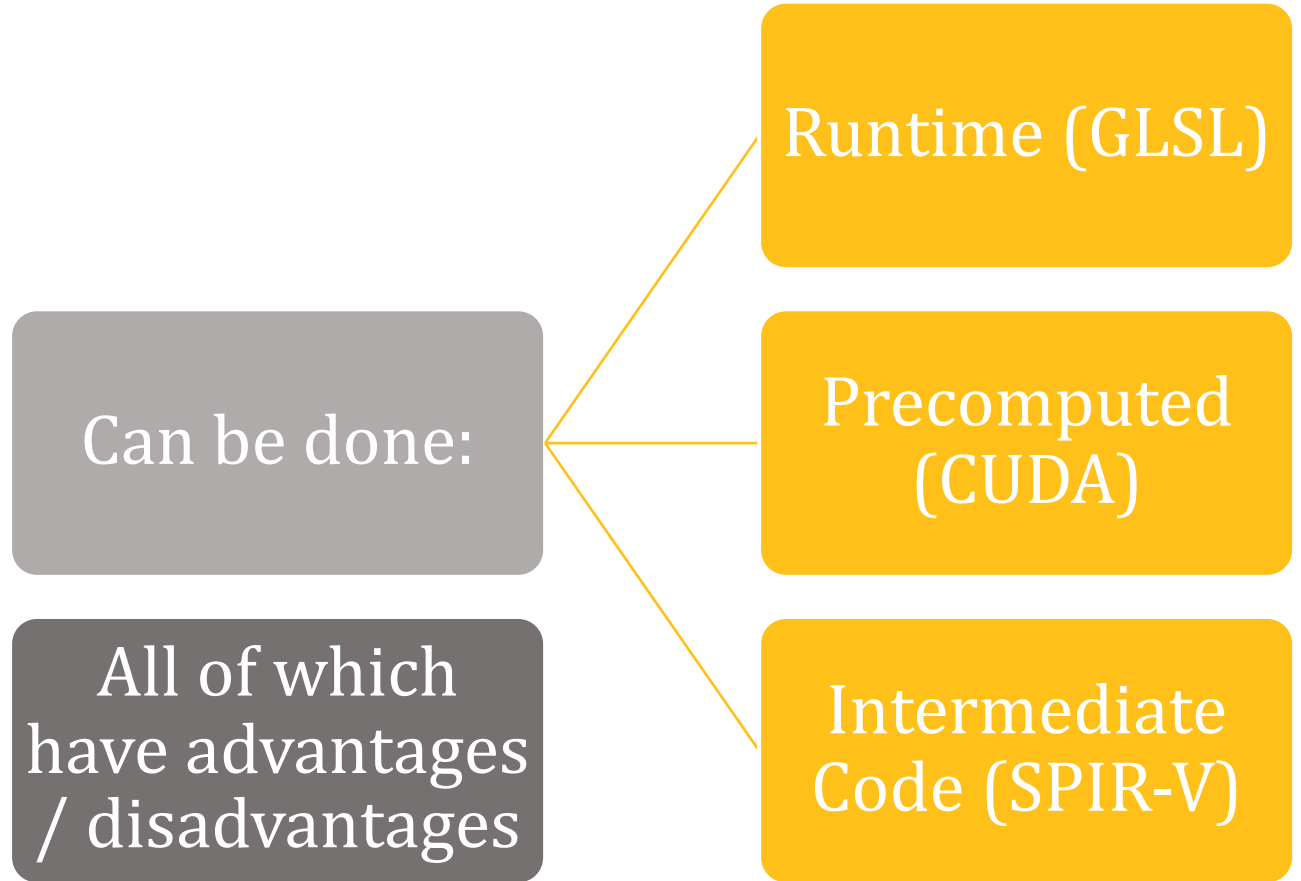


We need to discuss:

Compilation
Language
I/O



Shader Compilation



SPIR-V Compilation



Shaders precompiled to modules (libraries)

Collections of functions with entry points

Each has a name

And a type (which pipeline stage it is for)



Stored as a stream of 4B words

Essentially, an opcode / bytecode like Java

Can be inspected with spirv-dis



GLSL

- Essentially a dialect of C
- With some C++ conveniences
- We (like Vulkan) will use GLSL for this and next lecture
- Others are similar
- Most of the standard library routines built-in
 - Except memory allocation & I/O

GLSL Types

- `bool`: boolean type, as C++
- `int` / `uint`: basic integer type (usually 4B)
- `float` / `double`: IEEE floating point
 - `float` is often a lot faster than `double`
- `vec2`, `vec3`, `vec4`: floating point vectors
 - integers / doubles also available
- `mat2`, `mat3`, `mat4`: matrices




Matrix Types

- Float & double are available, no int or bool
- Represented in column-major format
- Can use `m[3]` to refer to column 3
- All usual arithmetic operations defined
- But you'll still need a library on CPU

I/O

- Shaders have no print routines
 - Makes it harder to debug
- Instead, they have shared buffers
 - Which change from time to time
 - So they need to be declared explicitly
- Based on *resources* and *descriptors*
 - Resource: a variable outside the shader
 - Descriptor: a bundle of resources





Descriptor Set

- Set of resources bound as a group
- Typically textures, samplers or buffers
- Set up as part of the pipeline
- Then bound to the inputs of the shader
- Push constants are a special case
 - Variables set *directly* from command buffer
 - All others are set by storage in a buffer

Shader Memory Access

Shaders can have local variables

These are usually in registers

Anything else is a resource (three kinds):

- Uniform Blocks (Read Only)
- Texel Buffers (Read Only)
- Shader Storage Blocks (Read/Write)

Read Only solves parallel problems

Uniform Blocks

- Read-only memory
- Shared between all invocations of the shader
- Limited size
- Usually the fastest memory (ie. cache)
- Declared with the uniform keyword



Uniform Buffers

- Vertex (shader storage) buffers hold attributes
 - Values that change for every vertex
 - Expensive to change (20K vertices => 2MB)
 - Should be loaded at startup/level load
- Uniform buffers are for constants
 - These are the control variables
 - Often small in size (< 1KB)
 - Can be modified every frame



Types of Uniform Data

Transformation matrices

- Small, change every frame

Flags & constants to control rendering

- Small, may change every frame

Textures

- Large (MB+), often constant
- Unless they are *generated* in multipass



Transformation Matrices

- We may want:
 - Model matrix
 - View matrix
 - Projection matrix
 - Texture matrix
 - Shadow matrix
 - And many others ...

Multiple Matrices

- OpenGL had matrix *stacks*
- Useful for animation
- But this broke down with skinning
 - Vertices affected by *multiple* transformations
 - E.g. armpit – skin affected by chest & arm
 - Need to compute an average position
 - So need access to more than one matrix

Vulkan[®] Solution

- Uniform buffer is untyped (like all buffers)
- Store anything you want
- You have to specify alignment / offsets
 - So shader can find the data it needs
- But you can pass anything you want
 - Including 50 different matrices



Uniform Buffer Stages

- Control code specifies a descriptor layout
 - As part of pipeline creation
 - Then allocates a descriptor set from a pool
 - And binds the descriptor set for rendering
- Layout (again) tells shader how to access it
- Very similar to vertex buffer setup
 - But don't use a staging buffer (why not?)



Image Buffers (Textures)

Textures are basically constants

- Shared between many vertices
- So they are passed as uniform buffers
- Ideally, no texture swapping during the pass

But they need an extra layer

- A *sampler*
- Which takes care of interpolation / filtering

Texel Buffers

- Read-only
- Can convert formats for you
- Can do interpolation
- Best choice for large arrays of data
- Accessed with `texelFetch()` function
- Declared as uniform `samplerBuffer`



Shader Storage Blocks

- Read-write access
- Support atomic operations
- Therefore often slower
- But much larger in size
- Declared with buffer keyword
- For example, a vertex buffer
- Or a frame buffer



A Simple Vertex Shader

```
// Taken from the Vulkan Tutorial
// These are needed to specify version, &c. to the compiler
#version 450
#extension GL_ARB_separate_shader_objects: enable

//////////////////////////////////////
//
// I/O VARIABLES: Set up through pipeline instantiation
//
//////////////////////////////////////

// This declares an output variable
// I'll leave it to you to work out exactly what the type information means
layout(location = 0) vec3 fragColor;

// Why haven't we declared an input variable?

//////////////////////////////////////
//
// LOCAL VARIABLES: instantiated for every invocation
//
//////////////////////////////////////

// Declare three vertex positions in 2D
vec2 positions[3] =
// And instantiate them
    vec2[] (
// Notice the explicit invocation of constructors
        vec2(0.0, -0.5),
        vec2(0.5, 0.5),
        vec2(-0.5, 0.5)
    );

// and the main routine
void main()
{ // main()
    // gl_Position is one of the default output variables
    // gl_VertexIndex is one of the default input variables
    gl_Position = vec4(positions[gl_VertexIndex], 0.0, 1.0);
    // fragColor we declared ourselves, but it's also output
    fragColor = colors[gl_VertexIndex];
} // main()
```



And a Fragment Shader

```
// Again, this is needed for the compiler
#version 450
#extension GL_ARB_separate_shader_objects: enable

// the fragment colour (input)
layout(location = 0) vec3 fragColor;

// the output colour
layout(location = 0) vec3 outColor;

// and the main routine
void main()
{ // main()
  // just copy the input colour to the output
  // note that interpolation has already happened
  // in the rasteriser
  // also notice the 1.0 set for the opacity
  outColor = vec4(fragColor, 1.0);
} // main()
```



Images by

- S5 Ben Lambert
- S6 Victor
- S7 Fakurian Design
- S8 Pascal Bernardon
- S9 Adrianna Geo
- S10Remy Gieling
- S11 Hello I'm Nik
- S21 Ryan Quintal
- S22 Greg Rosenke

