

13 - Filtering

Dr. Hamish Carr & Dr. Rafael Kuffner dos Anjos

Agenda

- Images
- Sampling
- Filtering
- FFT



What Is an Image?

- An *ideal* or *continuous* image is a function
 - $I(x, y): \mathbb{R}^2 \rightarrow \mathbb{R}$
 - That maps location on the image to intensity
 - We should also represent *wavelength*
 - $I(x, y, \lambda): \mathbb{R}^2 \times \mathbb{R} \rightarrow \mathbb{R}$
 - Or even irradiance:
 - $L(X, \omega, \lambda): \mathbb{R}^3 \times S^2 \times \mathbb{R} \rightarrow \mathbb{R}$



On the Retina

- Consider the irradiance on the retina
 - $I_{retina}(x, y, \lambda) = \int_{S^2} L(X, \omega, \lambda) dS$
 - Integrates all light at that wavelength
 - Over all incoming directions
- If lens is perfect, only one direction counts
 - But the retina is *not* continuous
 - It is made of cells of varying area



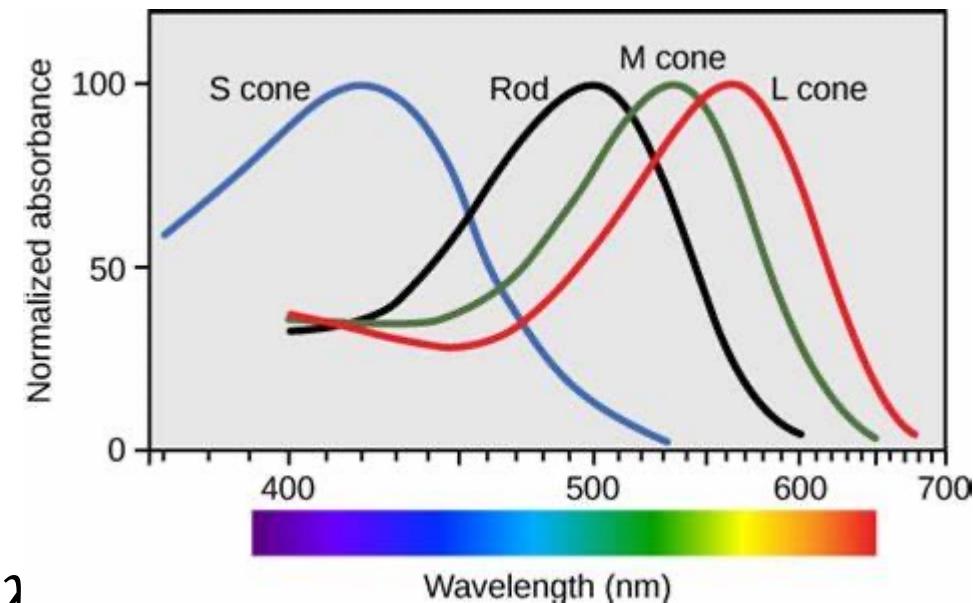
Discrete Retina / Image

- Retinal cells are blobs
- We will assume a nice square array
 - i.e. the pixels on the image plane
 - indexed on i, j
- So we can transform this to:
 - $p_{i,j} = \int_{x_i - \frac{1}{2}\Delta x}^{x_i + \frac{1}{2}\Delta x} \int_{y_j - \frac{1}{2}\Delta y}^{y_j + \frac{1}{2}\Delta y} \int_{S^2} L(X, \omega, \lambda) dS dy dx$



Colour Response

- Each cell has its own response function
 - $r(\lambda): \mathbb{R} \rightarrow \mathbb{R}$, $g(\lambda): \mathbb{R} \rightarrow \mathbb{R}$, or $b(\lambda): \mathbb{R} \rightarrow \mathbb{R}$
- How strongly it responds to wavelength λ
- It may actually vary between cells
- So assume we have $r_{i,j}(\lambda)$:
- $p_{i,j} = \int_{x_i - \frac{1}{2}\Delta x}^{x_i + \frac{1}{2}\Delta x} \int_{y_j - \frac{1}{2}\Delta y}^{y_j + \frac{1}{2}\Delta y} \int_{350nm}^{750nm} \int_{S^2} r_{i,j}(\lambda) L(X, \omega, \lambda) d\lambda dS dy dx$
- We could even parameterise with time t





Assumptions

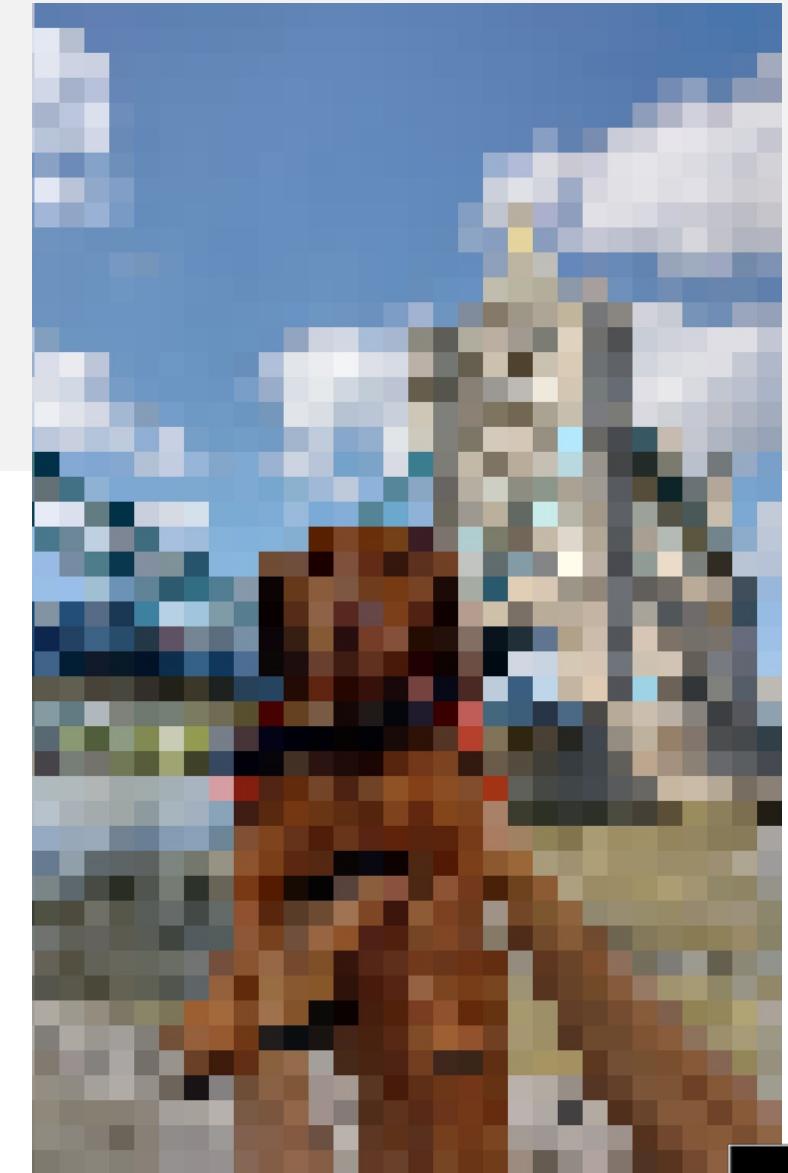
- Assume a perfect lens (only one ω matters)
- Assume square pixels (retinal cells)
- Assume one wavelength
 - $p_{i,j} = \int_{x_i - \frac{1}{2}\Delta x}^{x_i + \frac{1}{2}\Delta x} \int_{y_j - \frac{1}{2}\Delta y}^{y_j + \frac{1}{2}\Delta y} I(x, y) dy dx$
 - And assume we are storing as floats
 - not integers





Sampling

- We never actually compute this
 - although perhaps we should
- Instead, we sample the function
 - measure the light at the centre of the pixel
- Then display the same value across the pixel
- And this leads immediately to errors
 - In particular, *aliasing*



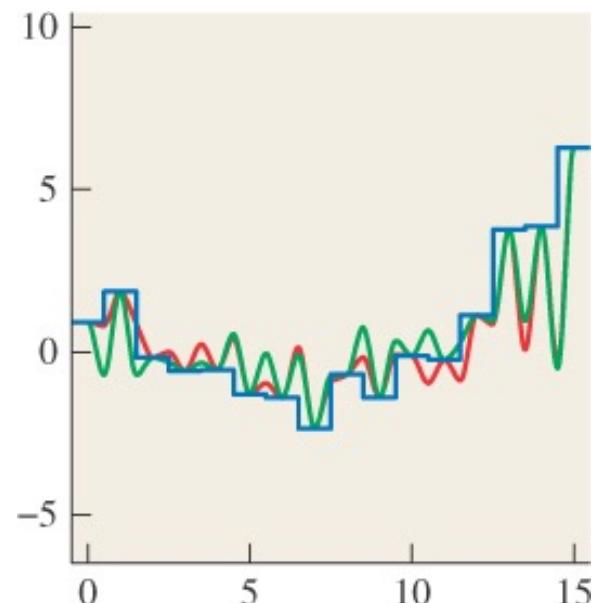
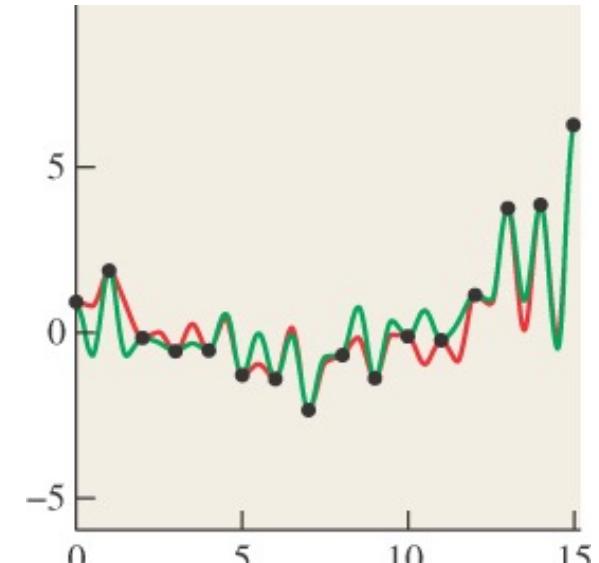
Reconstruction

- Suppose we measure a function at 10 points
- Can we *reconstruct* the function?
 - i.e. recover the continuous version
- Only if there is a unique representation
- And it is easy to show there isn't



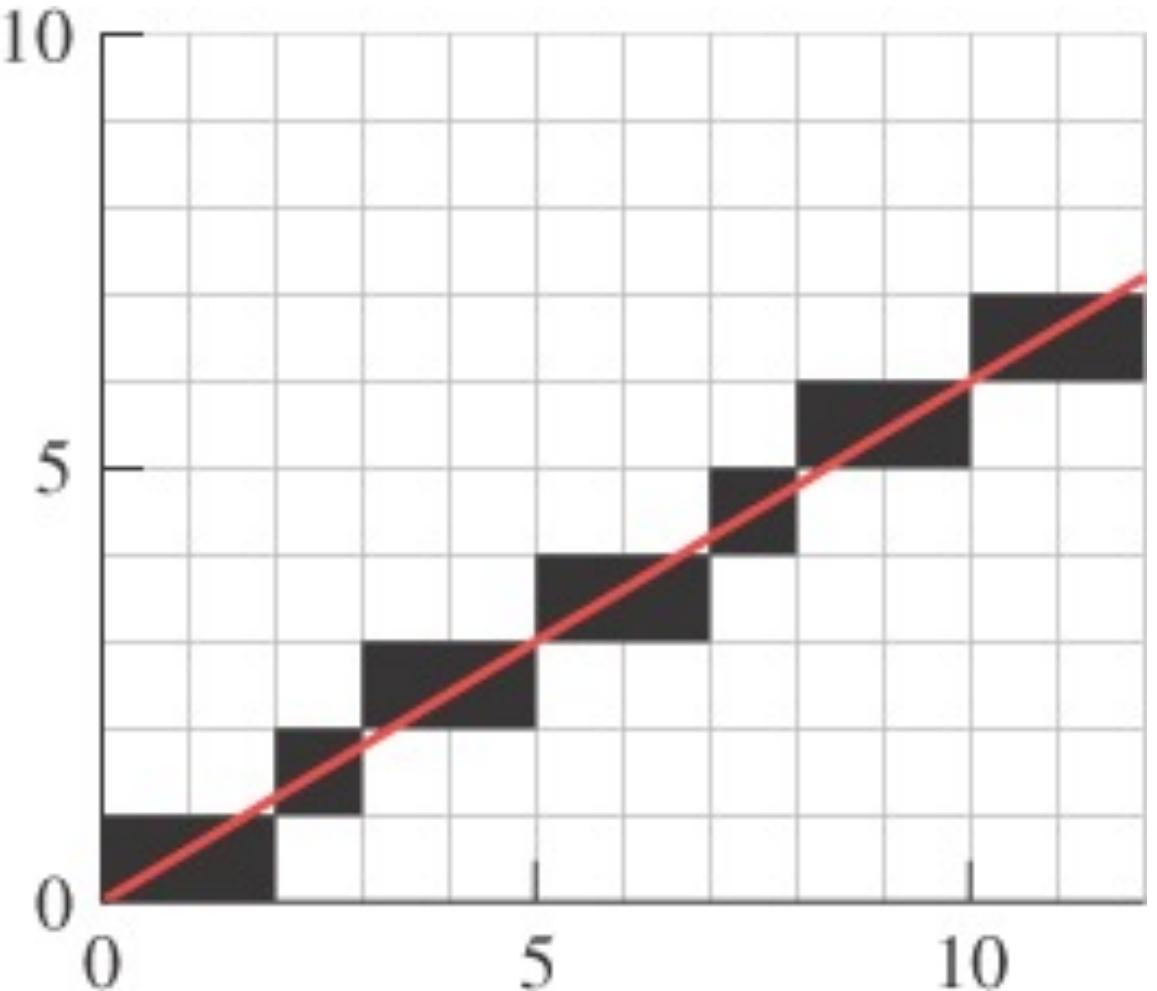
1-D Example

- Both green and red have the same samples
- Which one is right?
- The screen does a third thing
 - Because of hardware

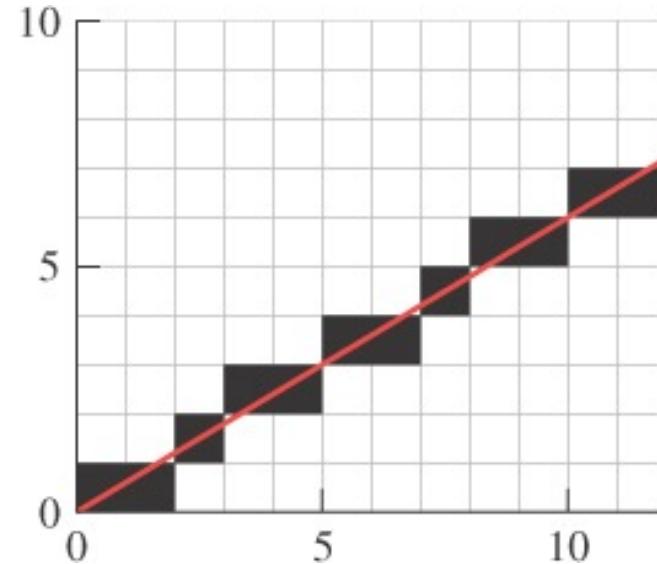


Aliasing

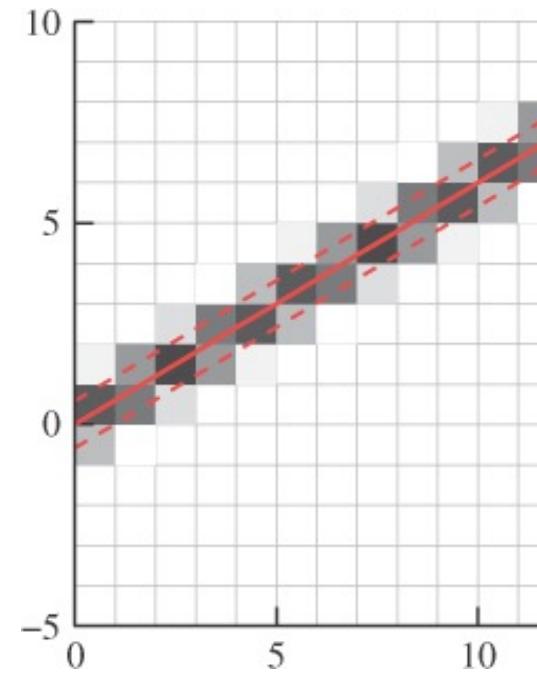
- A general phenomenon
- Caused by not enough samples
- Reconstruction dominated by error
- We *know* this should be a straight line
- But it is jaggy when discretised



Solution (Goal)



Aliased



Anti-aliased



Far away

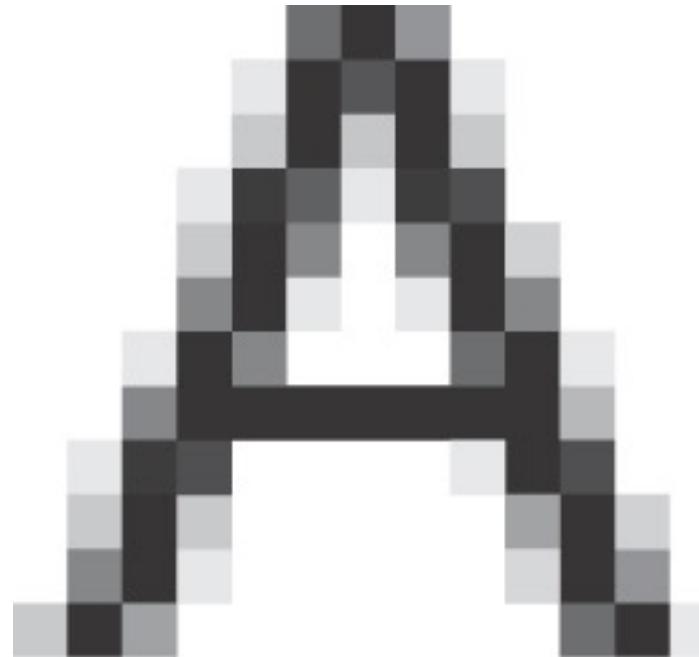
- Replace it with a grey-level
- Overlap between line of width 1 & pixel



Immediate Use: Fonts

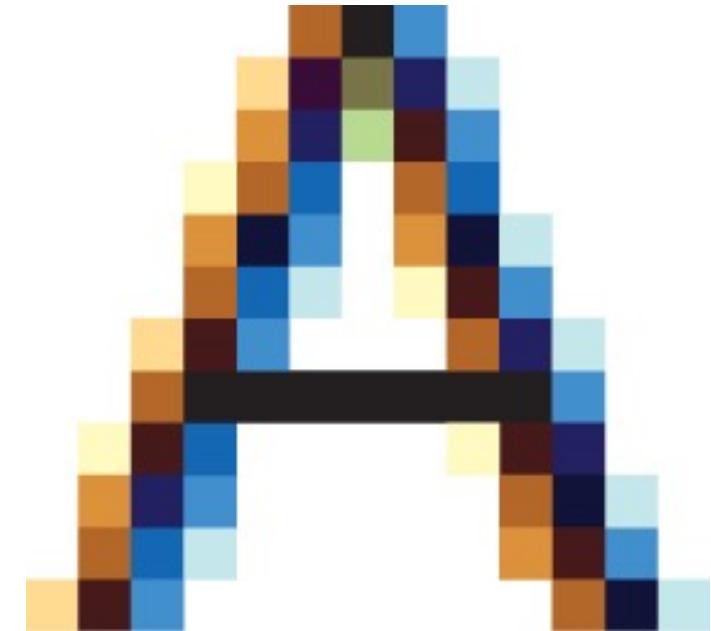


Binary Font



Anti-aliased

Using Greyscale



TrueType

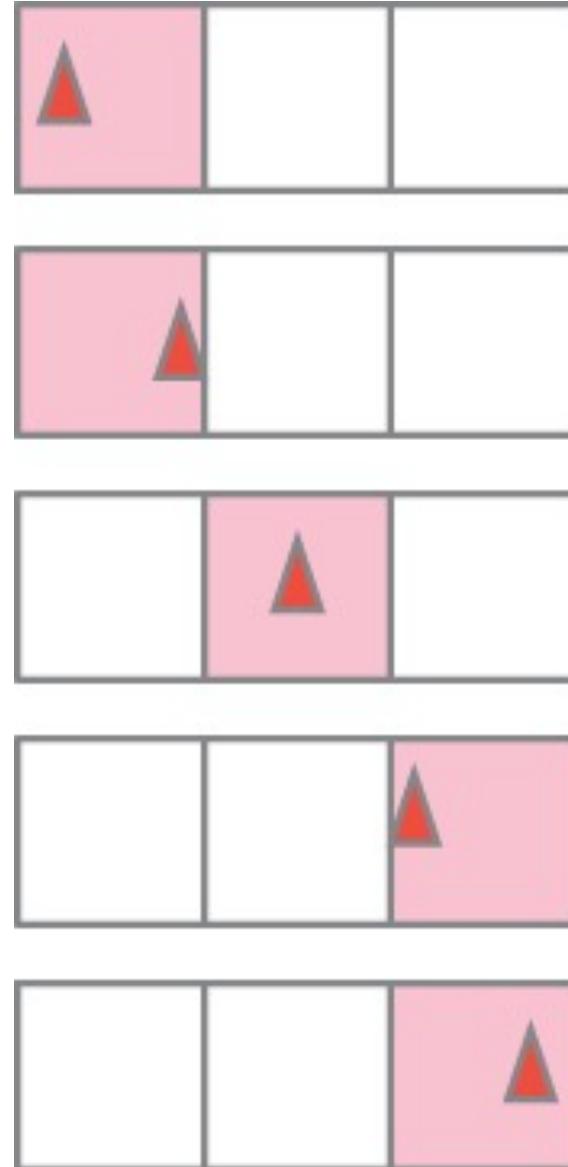
(RGB in different columns, eg LCD)



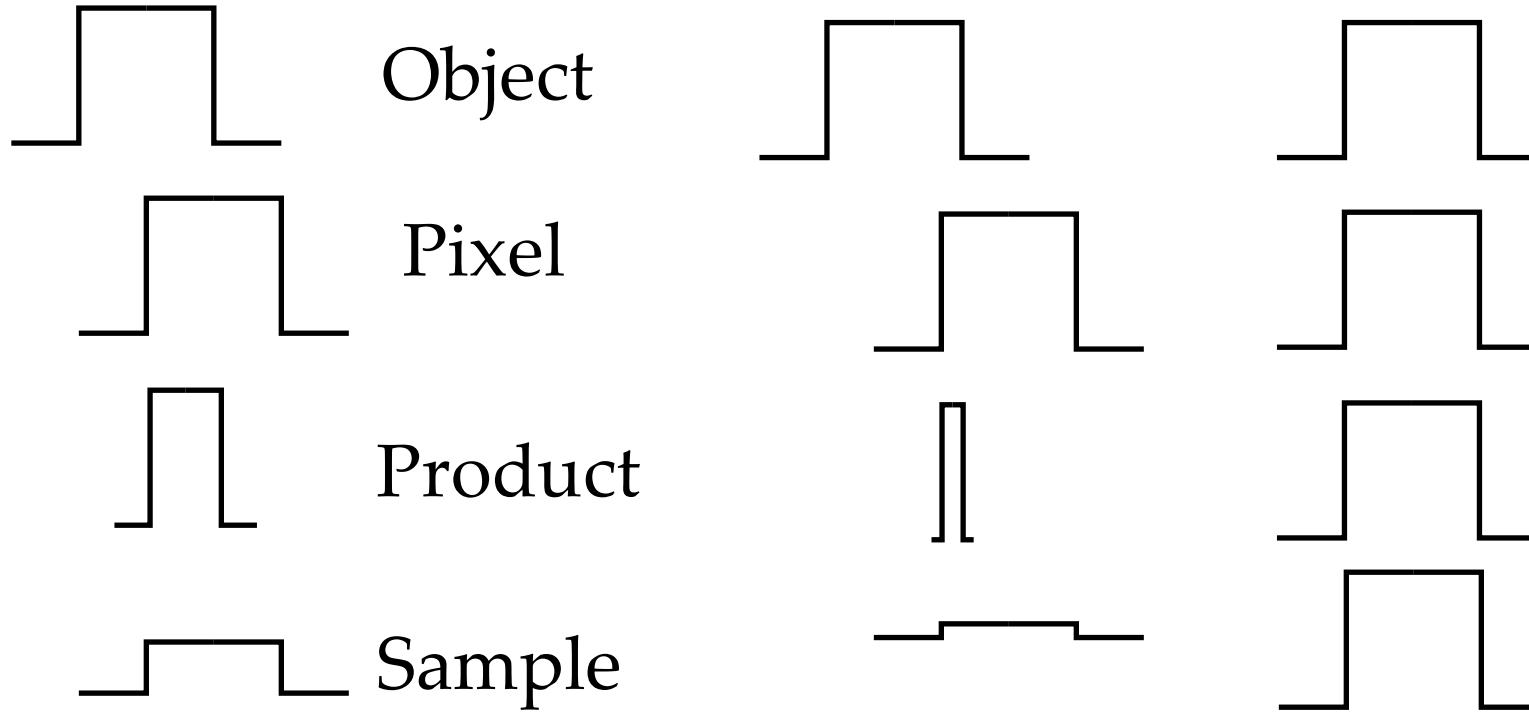
UNIVERSITY OF LEEDS

Aliased Motion

- Object moves 2 pixels in 4 time steps
- But it looks like it speeds up in the middle
- Really just the same problem



Solution



- Treat the object as a function itself
- Combine the function with the pixel

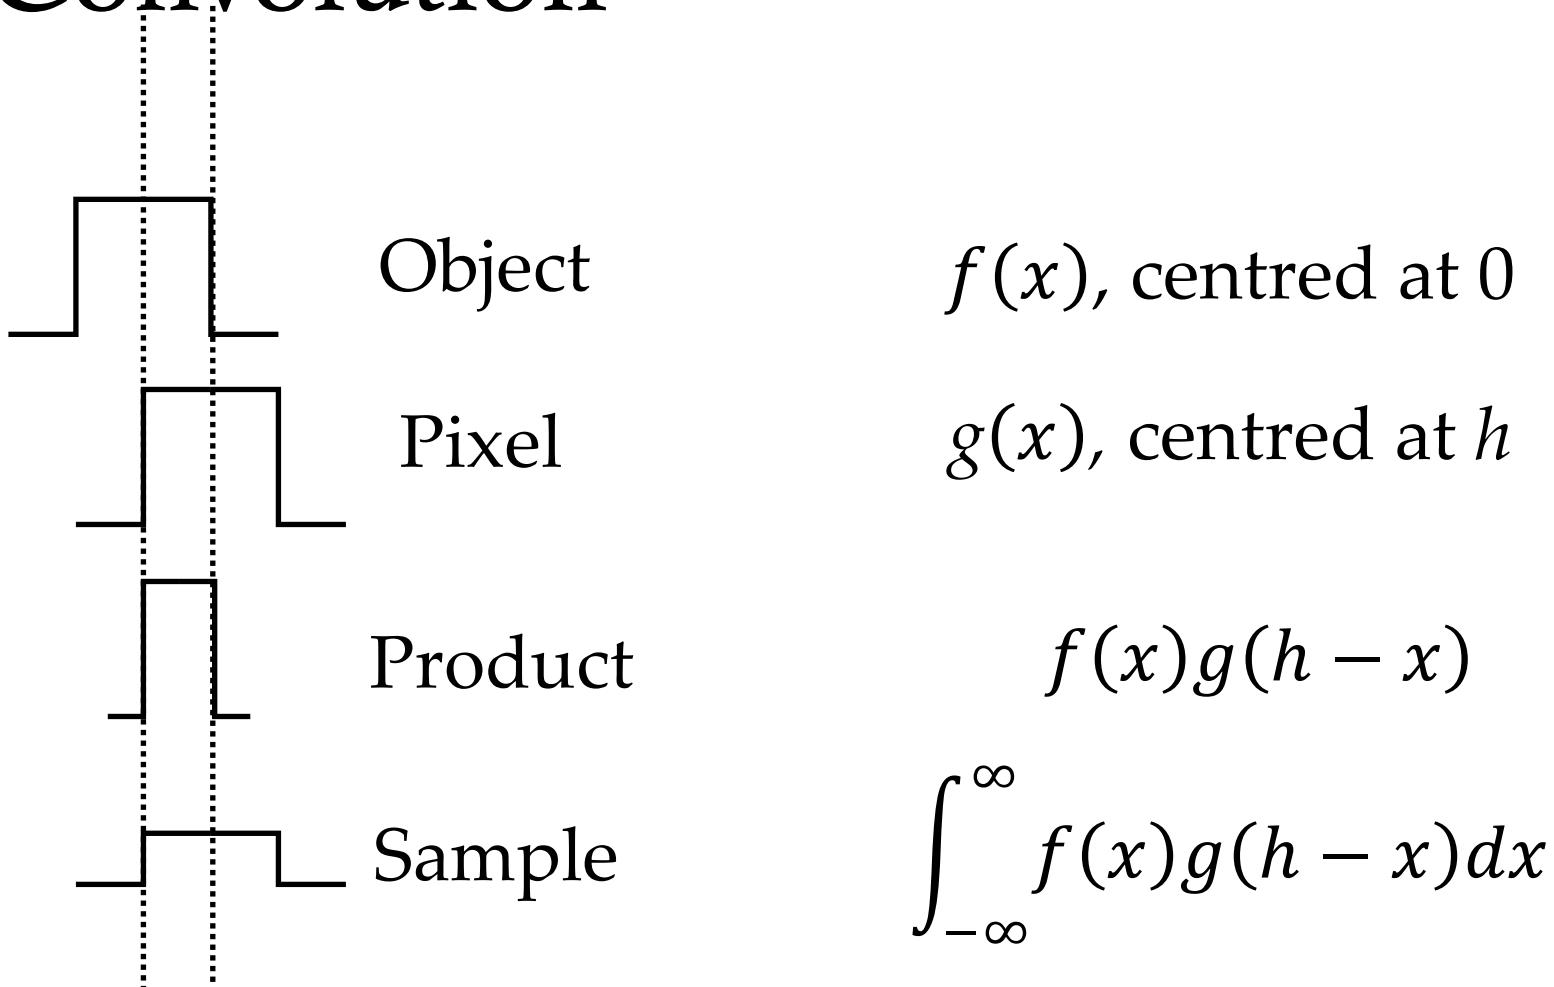


Mathematically

- Clearly, this varies with the *offset*
 - between object and pixel
- We express this as a function (of course)
- And an operation called *convolution*



Convolution

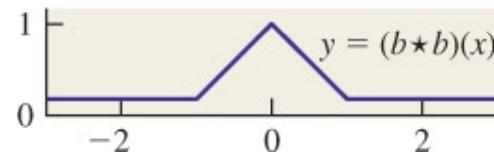


- Convolution works with any functions
 - and has nice mathematical properties

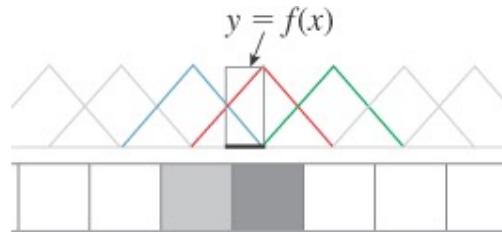


Filtering

- Depends on distance from pixel to object



- So the overlap varies as the pixel moves



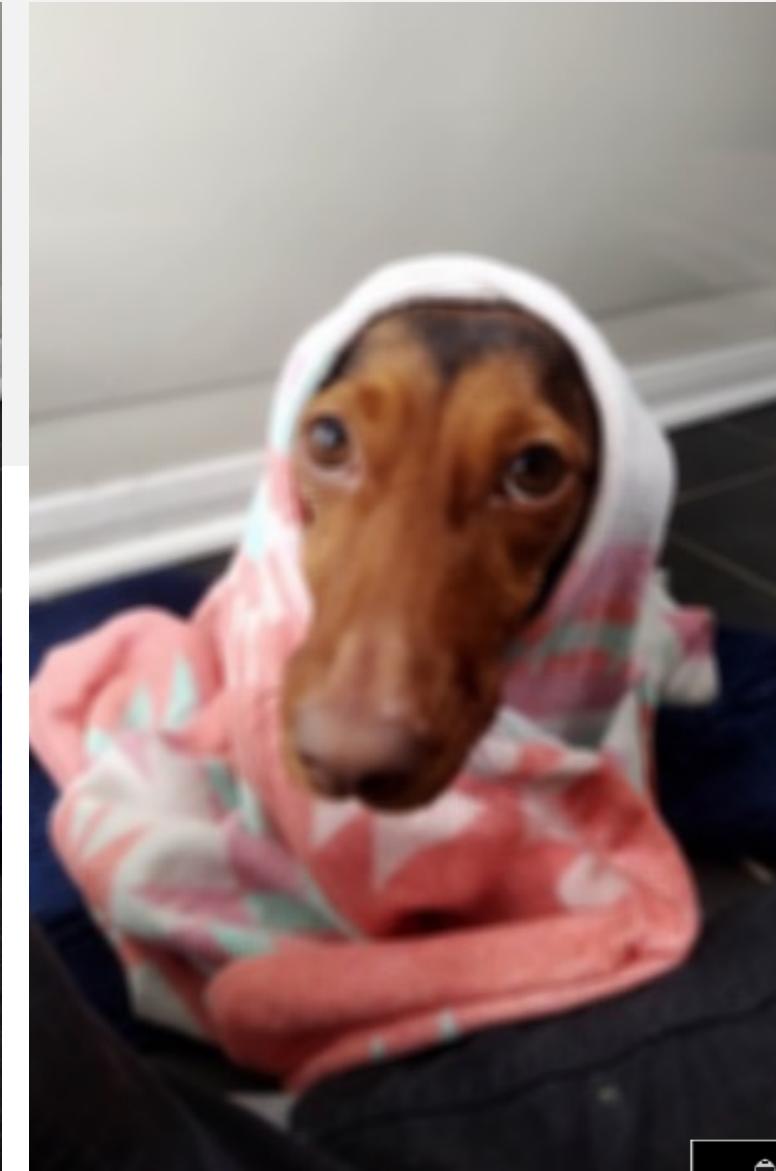
- But cost over an image is $O(NM)$
 - N is the number of pixels
 - M is the size of the filter





Blurring

- Simple trick: filter with a slightly larger area
 - average all adjacent pixels
- But this is just a 3x3 filter!
- $\begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$
- Larger filter, more blur.

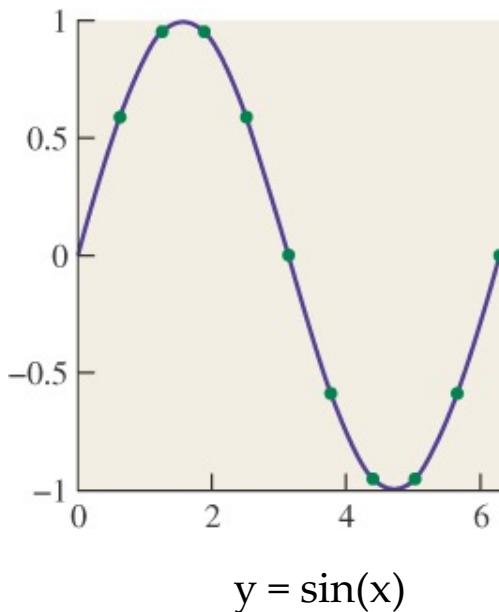
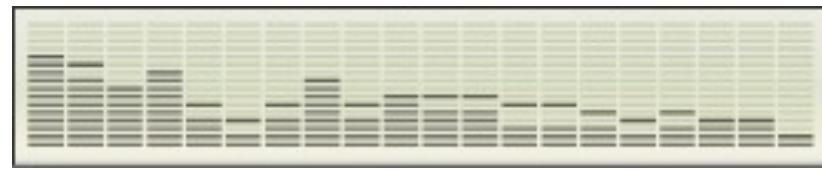
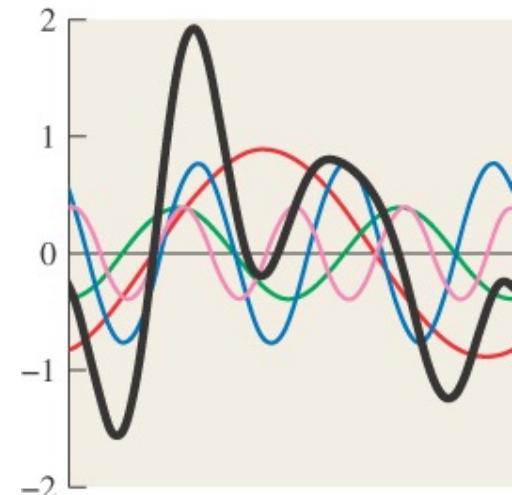


Sampling Theory

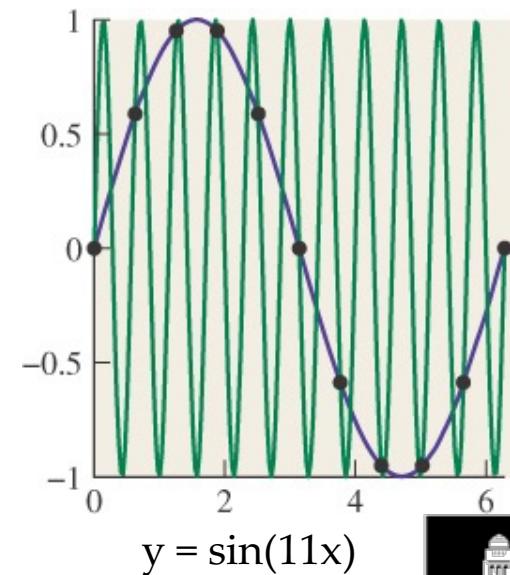
- The next piece of the puzzle
- Treats the image like an electrical signal
- Or an audio signal
- Breaks it down into *frequencies*
 - i.e. how rapidly it changes
- Expressed with Fourier analysis

Fourier Analysis

- Any *periodic* function can be broken down
- Into a sum of sines & cosines
- In the complex plane
- Based on our samples
- We will have to assume a *band limit*
- I.e. a maximum spatial frequency



$$y = \sin(x)$$

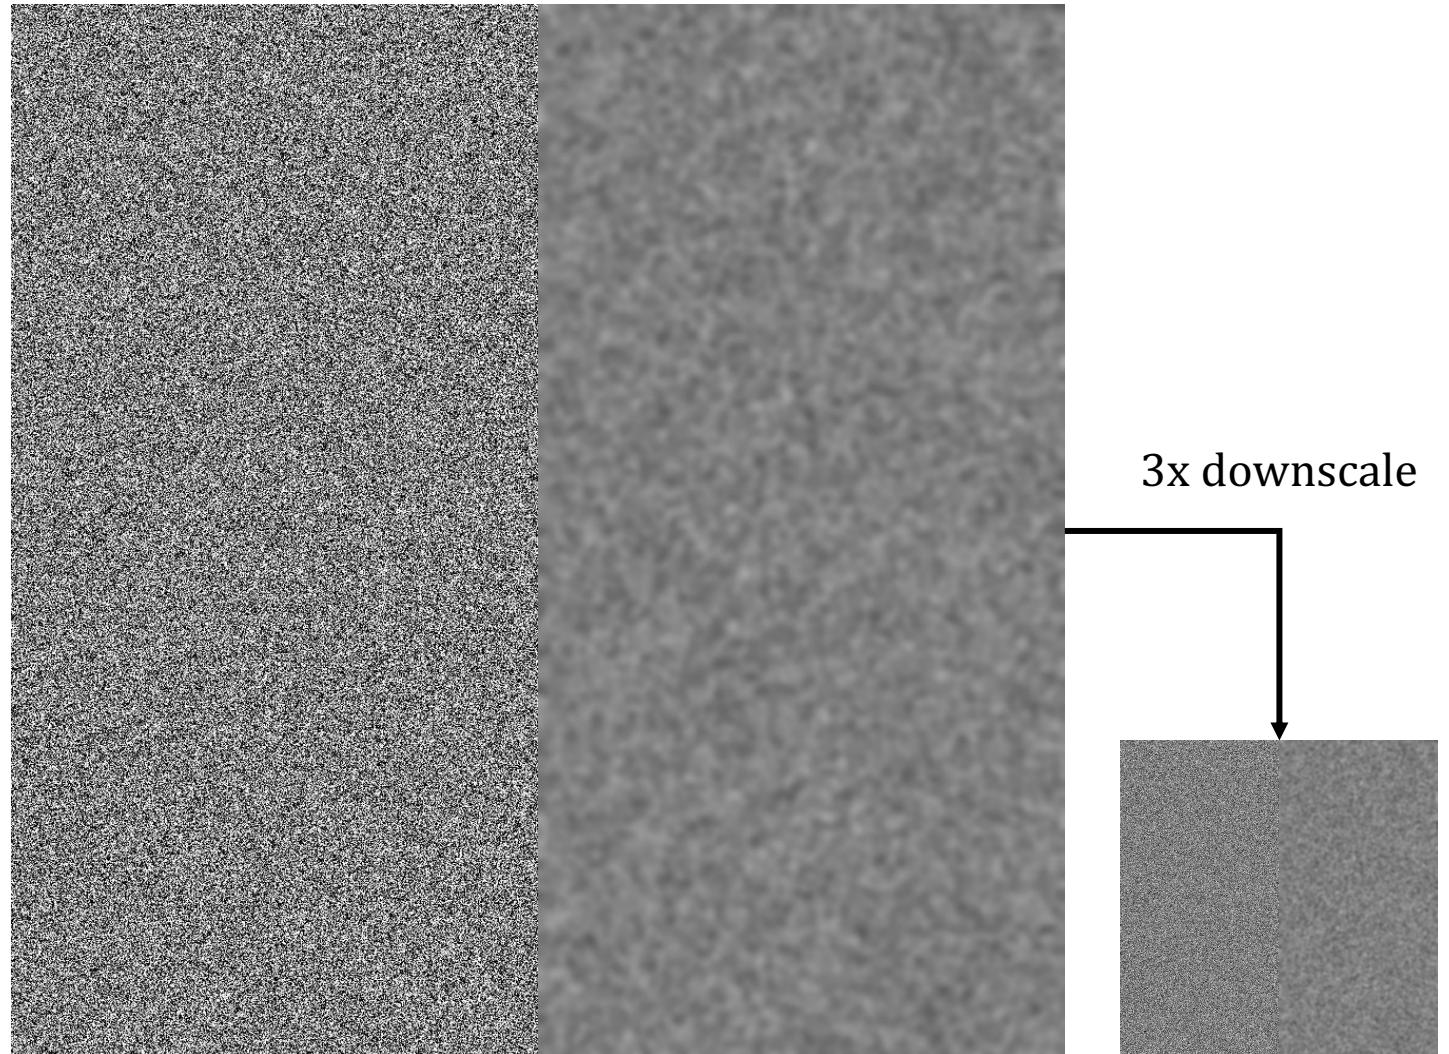


$$y = \sin(11x)$$



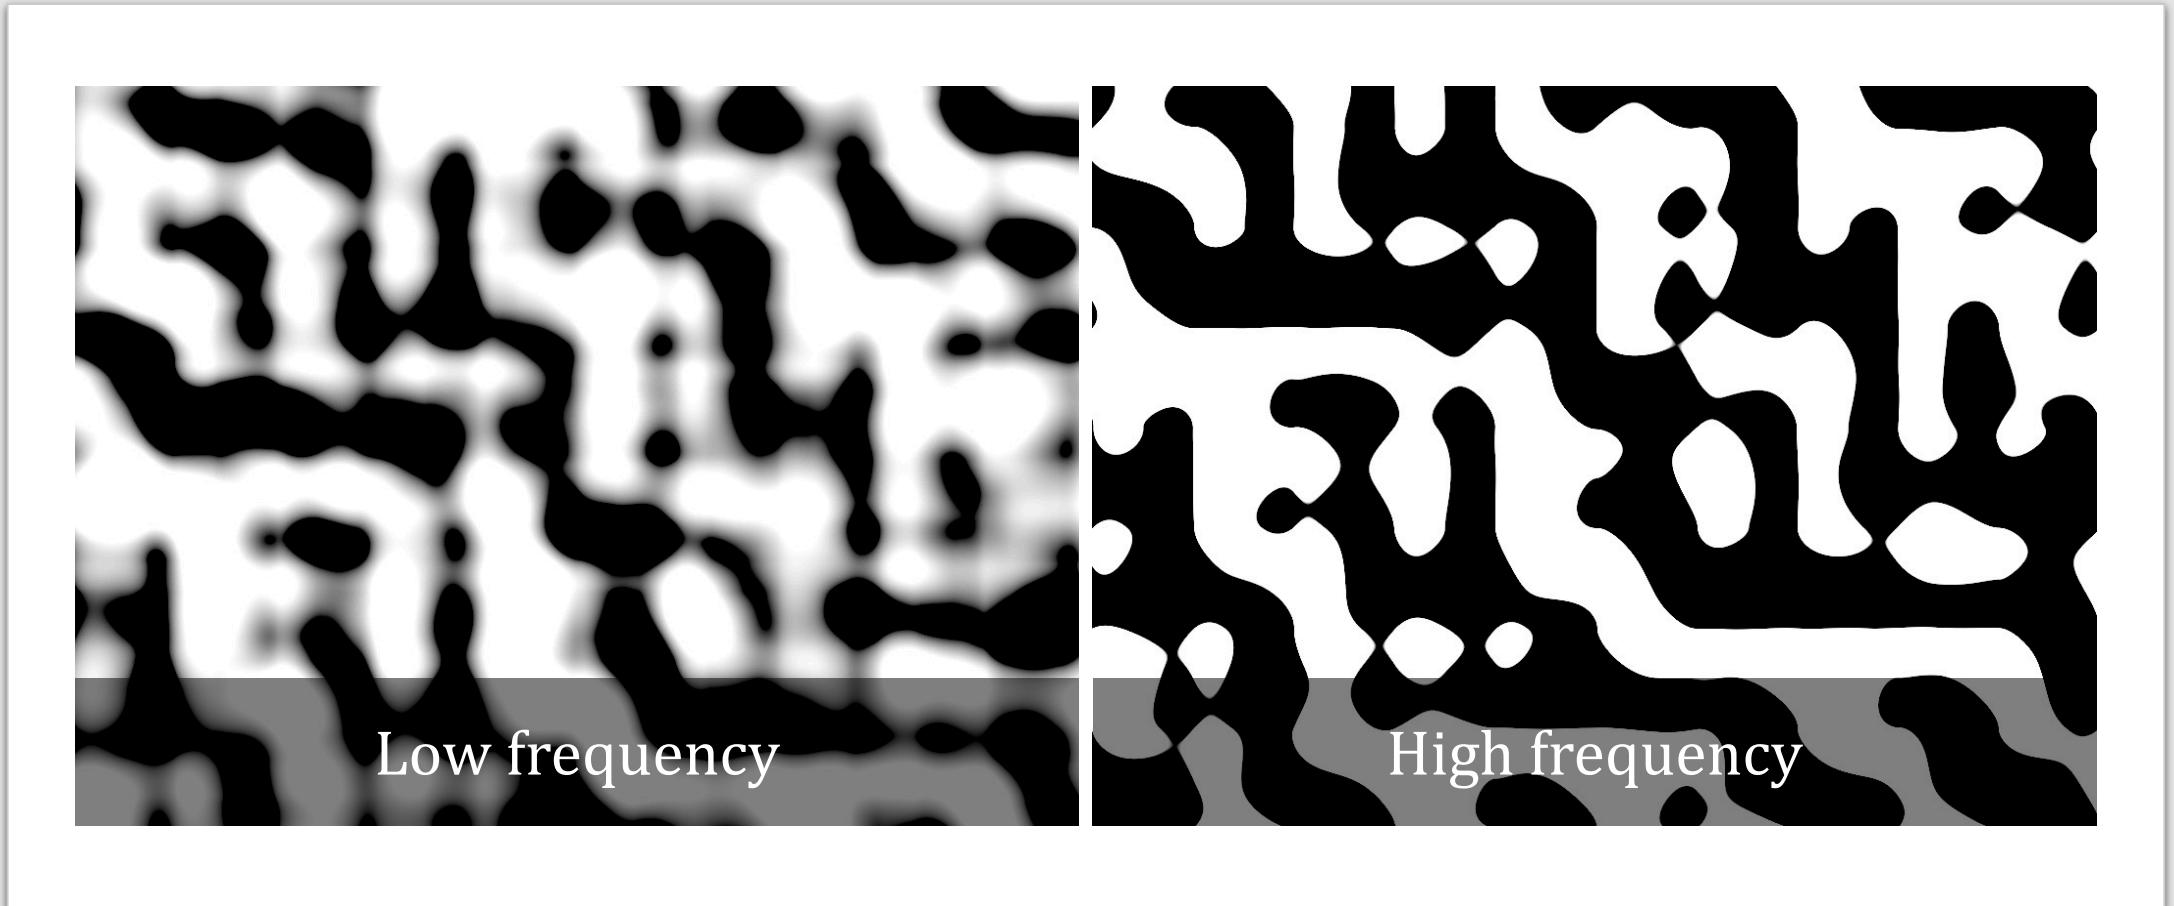
How can we see that in images?

- Each pixel is a sample of your signal
- Low resolution images:
 - band limited to low frequencies
- High resolution images:
 - Low frequency changes happen over several pixels
 - High frequency changes happen over fewer pixels



Seems uniform!





Example (Perlin noise)

Fourier Transform

- Rewrites our function
 - from a function of x
 - to a function of ξ : spatial frequency
- Expressed as samples:

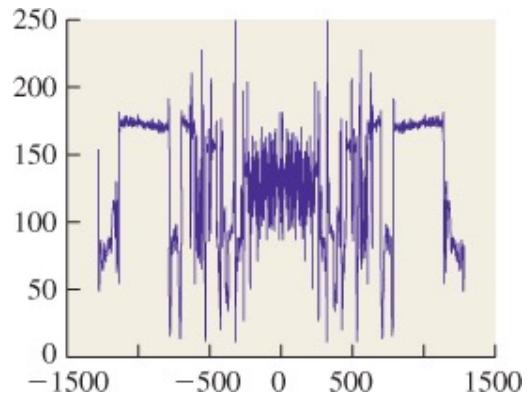
$$a_k = \int_{-\frac{1}{2}}^{\frac{1}{2}} f(x) \cos(kx) dx \text{ and} \quad (18.40)$$

$$b_k = \int_{-\frac{1}{2}}^{\frac{1}{2}} f(x) \sin(-kx) dx. \quad (18.41)$$

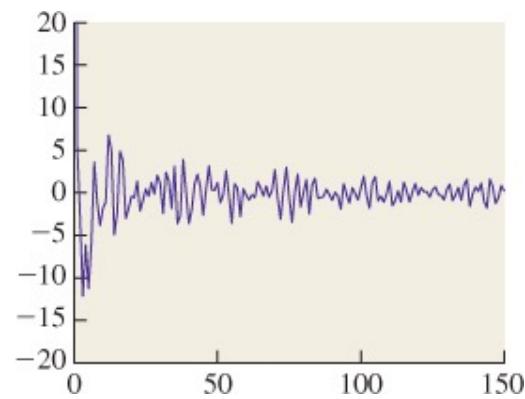
- actually a complex number $a_k + b_k i$



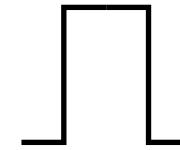
Frequency Filtering: 1D



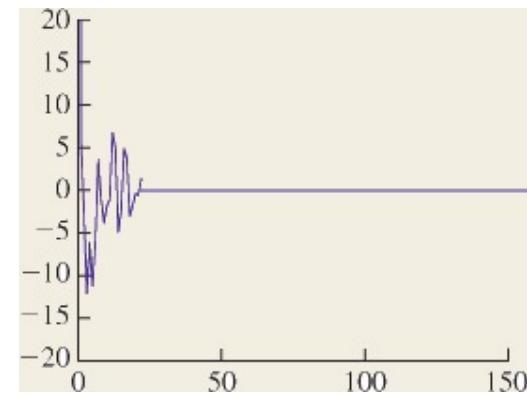
One Row



Fourier Transform



Box Filter

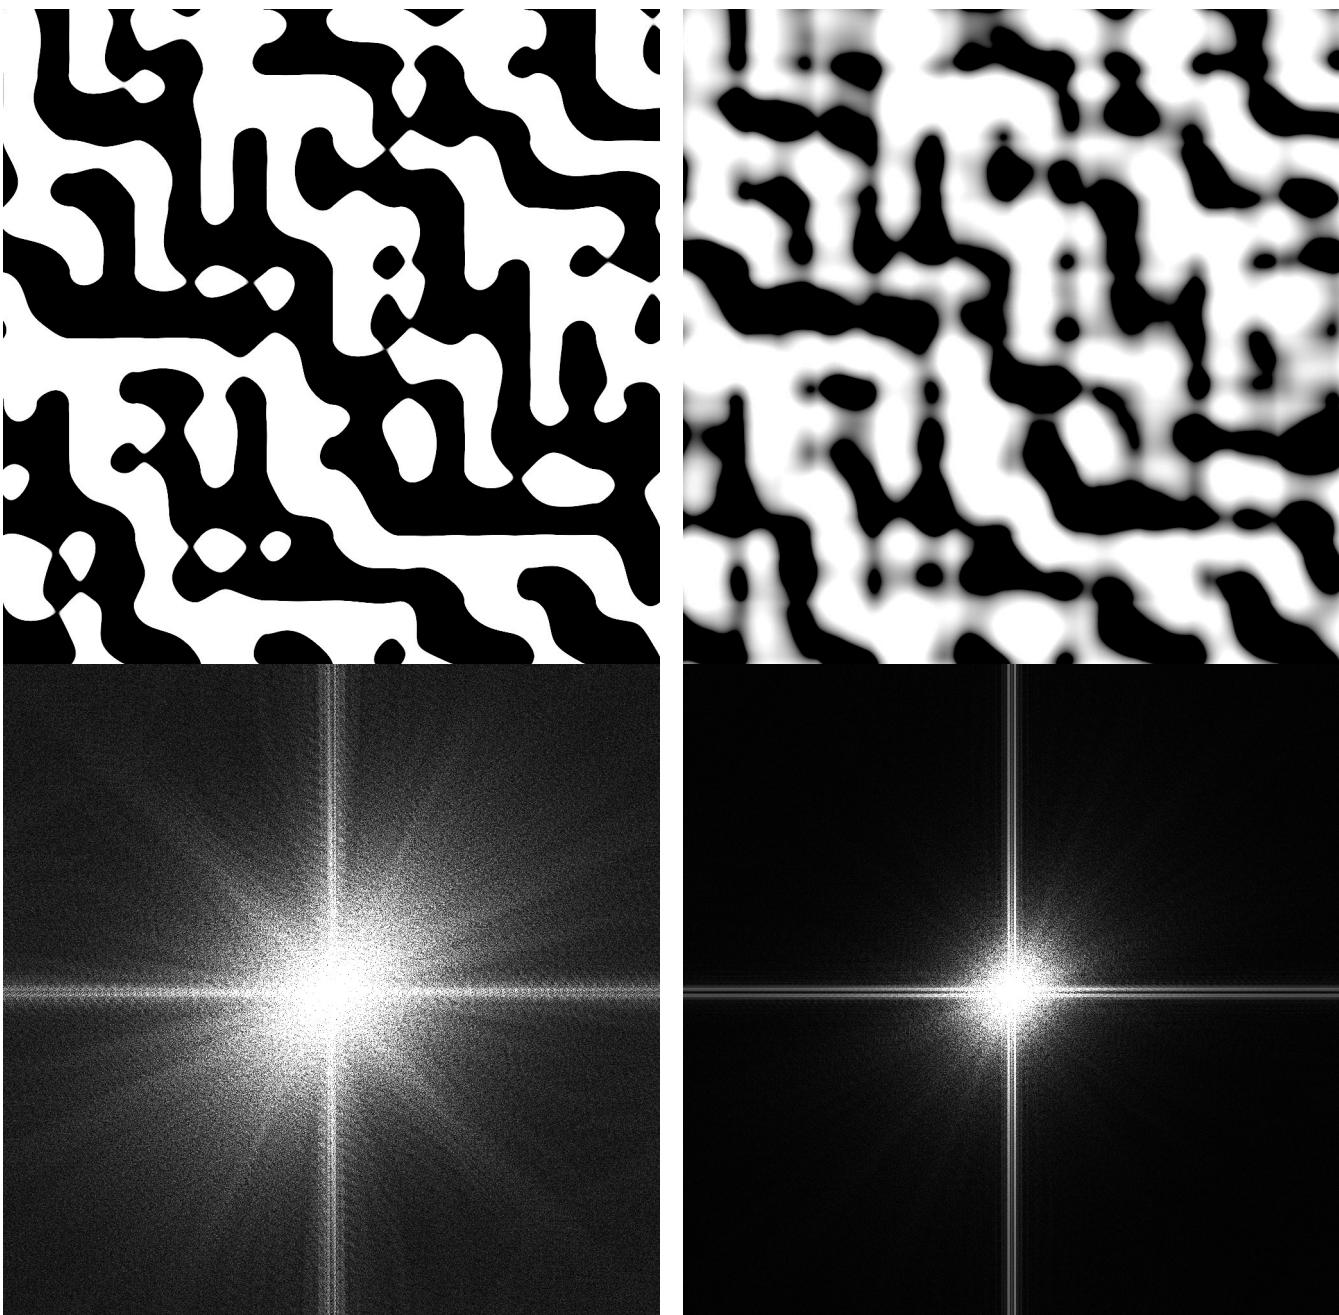


After Filtering



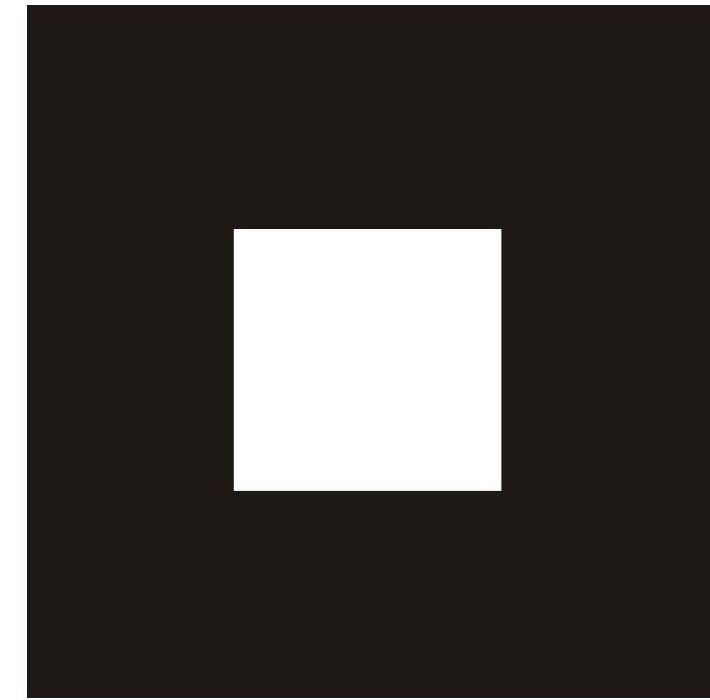
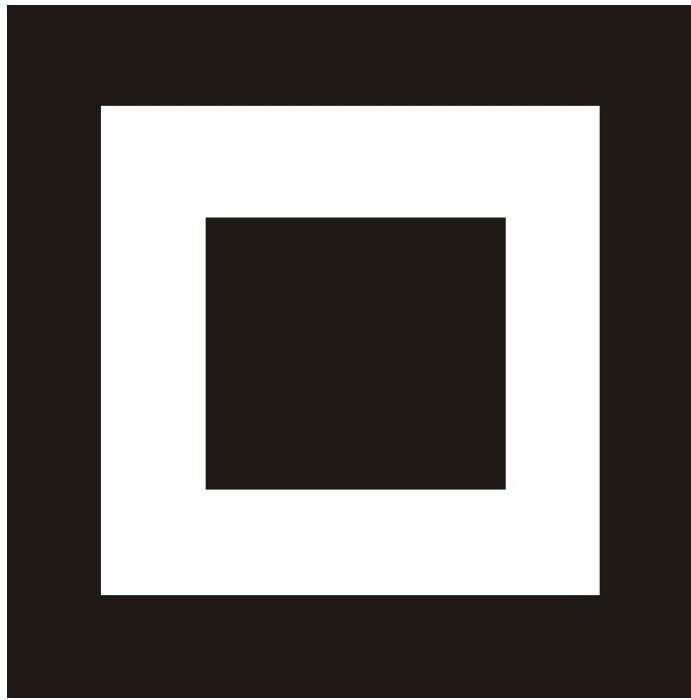
2D FFT

- We can plot the 2D fft to an image.
 - Typically shifted to the center
 - Lower frequencies in the center
 - Higher frequencies going towards the edges
 - Oriented lines mean frequencies of that orientation



Filtering 2D

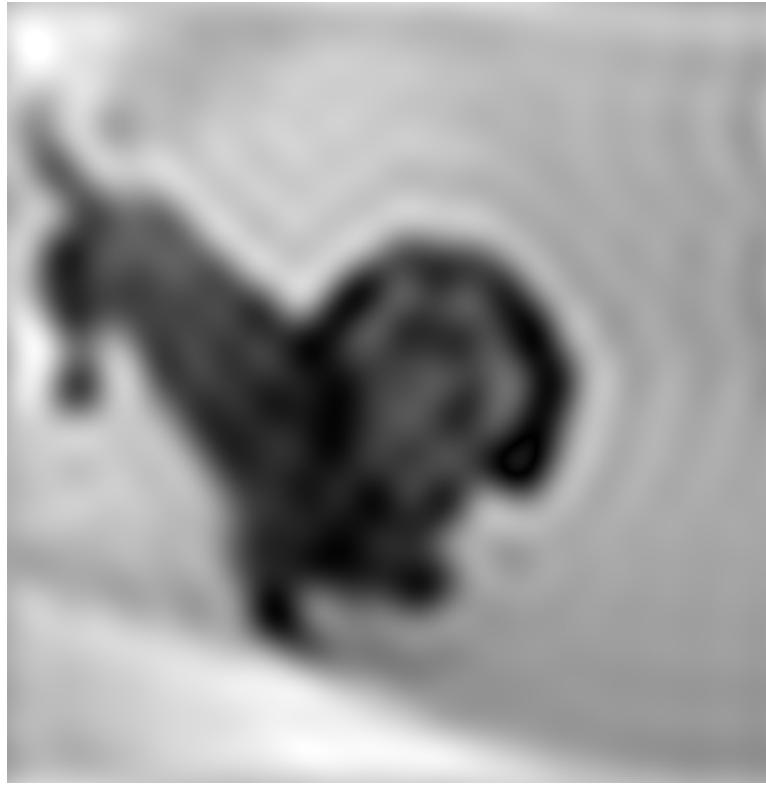
- Filters would look like this
- Can easily separate oriented frequencies as well
- Convolution = multiplication in Fourier domain.



Results

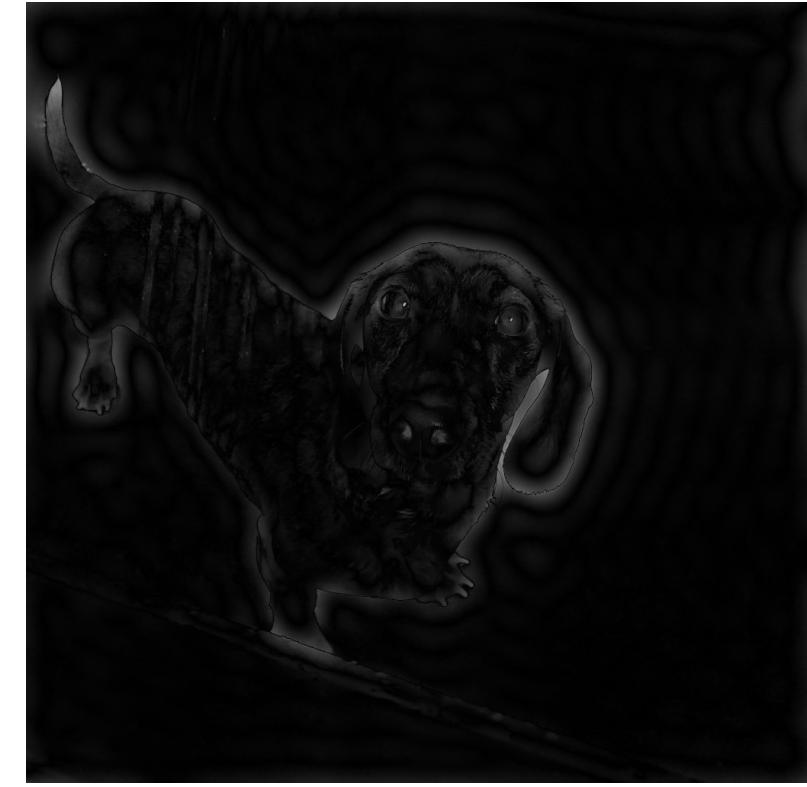


Lola



Low frequency

Ringing artifacts!

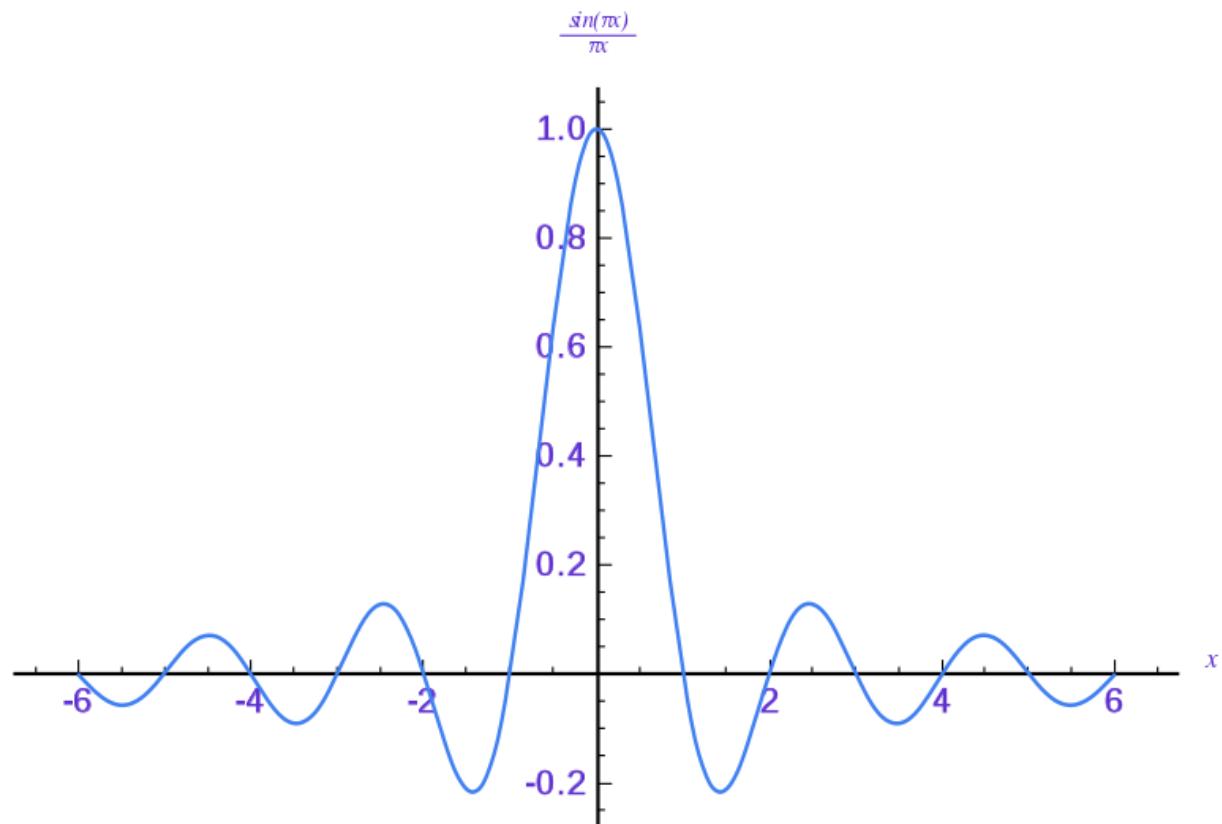


High frequency

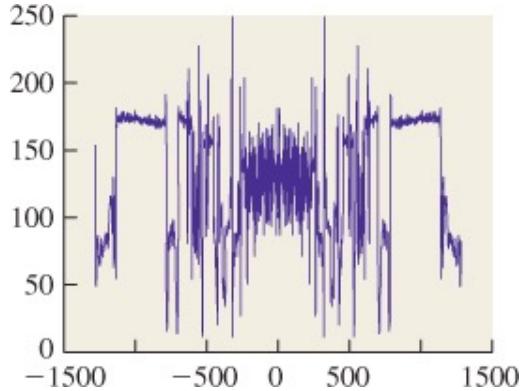


Sinc Filter

- What does this filter look like?
- In the spatial domain
- It becomes the *sinc filter*
- Ideal filter that removes frequencies above a certain cutoff, without affecting lower ones.

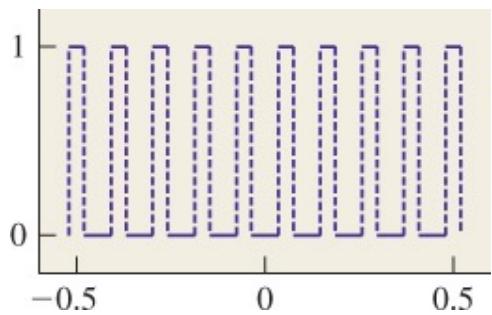


Sampling

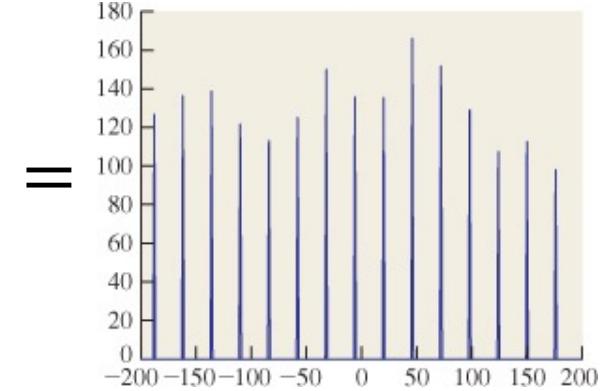


Taj Mahal

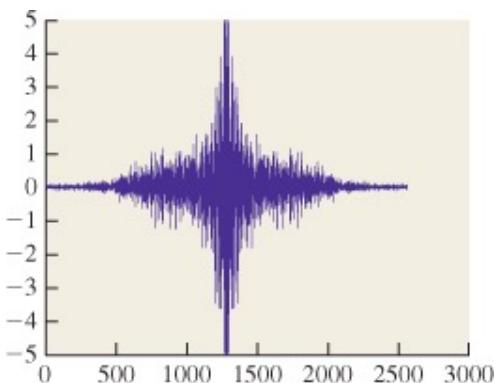
\times



Sampling
Function

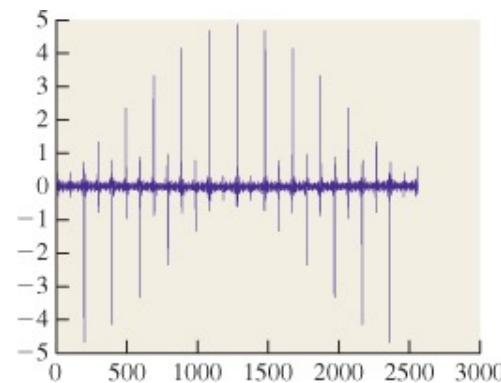


Sampled
Function



Fourier

Transform



Fourier

Transform

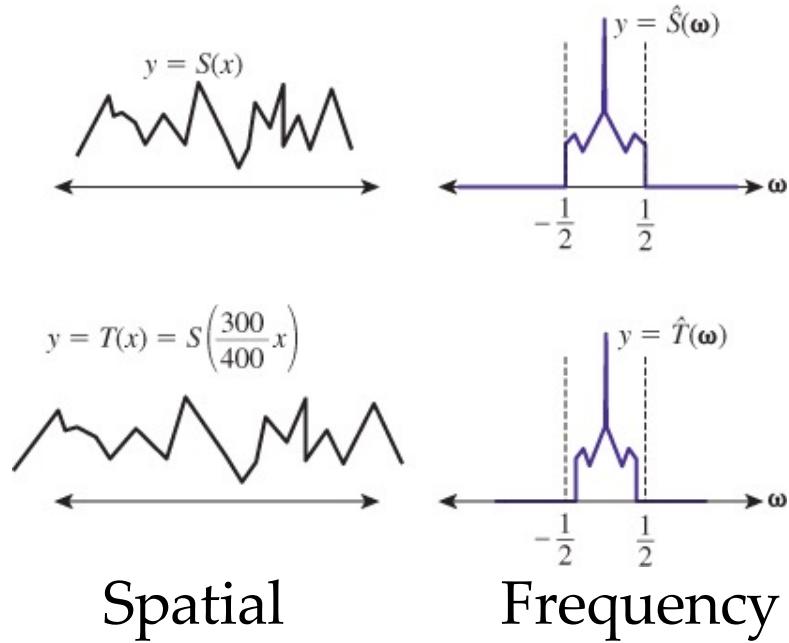


More generally

- Mipmapping is a special case
- We need a more general solution
- But scaling up & scaling down are different
 - Scaling up – use the sinc filter if possible
 - Scaling down – increases the frequency
 - And no longer band-limited



Scaling Up



Original Signal

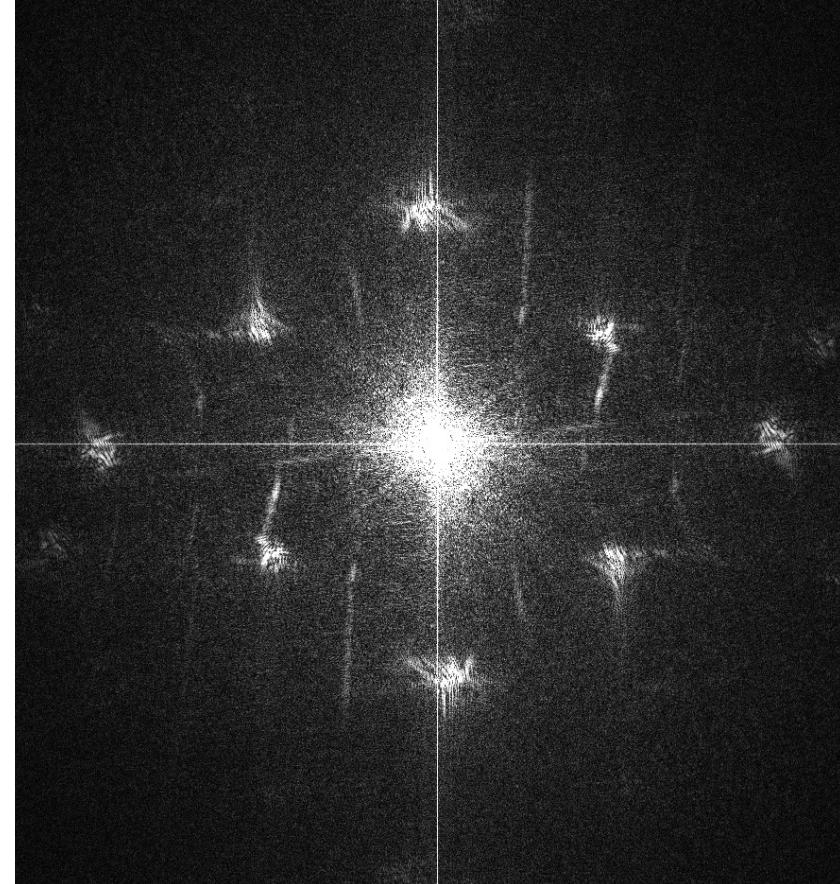
Rescaled Signal

- Increasing spatial resolution
- Decreasing frequency resolution
- Bandwidth limit is preserved

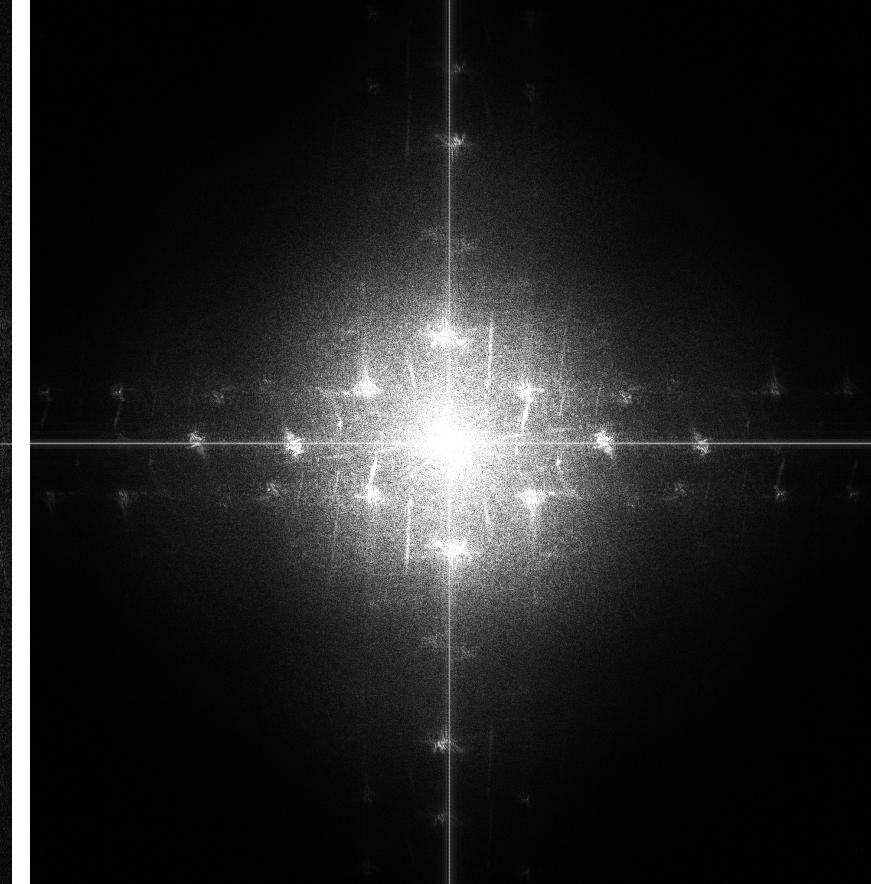




Image



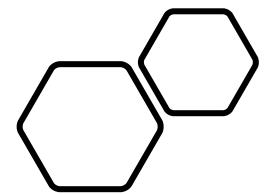
FFT

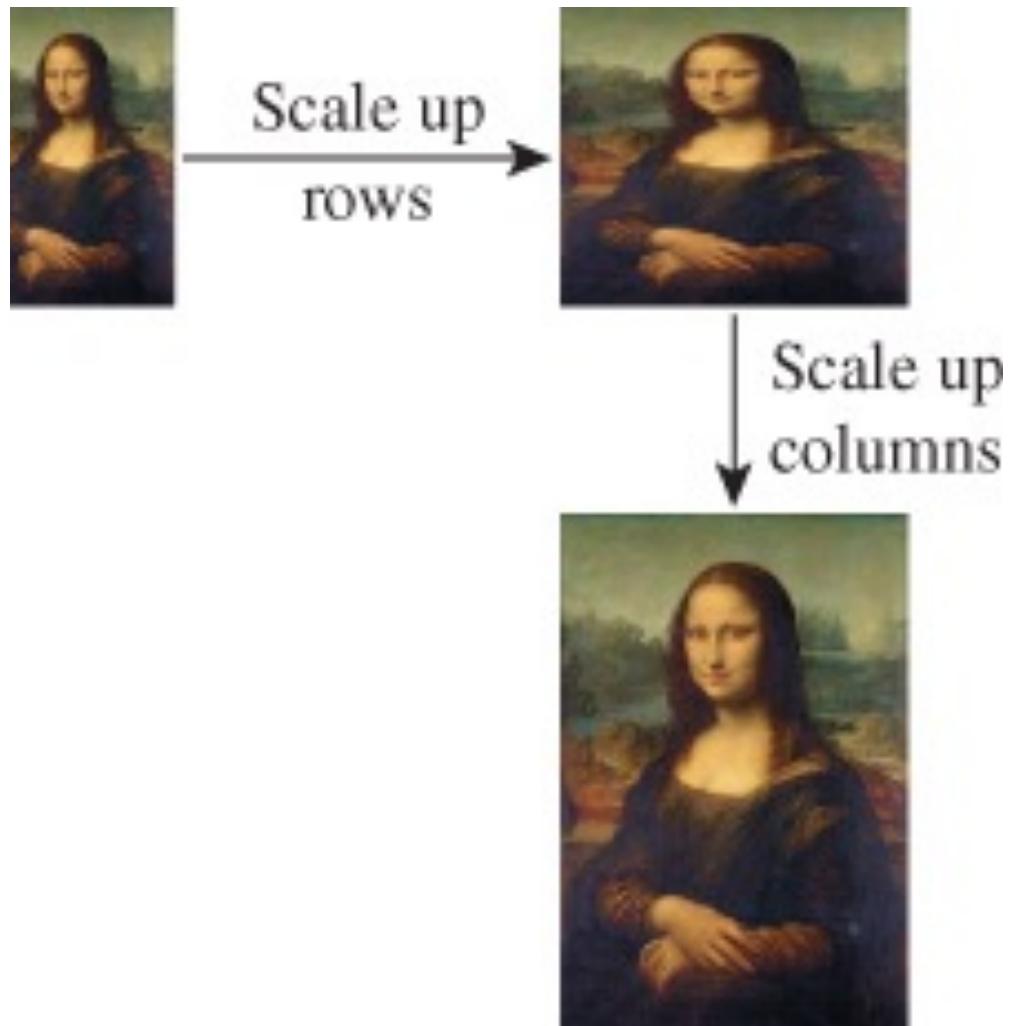


FFT of upscale

Example of
upscale

- FFT is “shrinked”



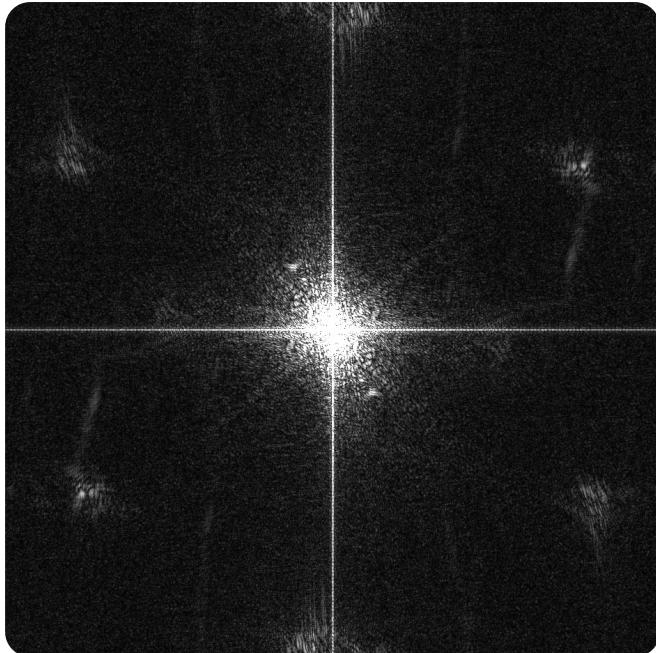
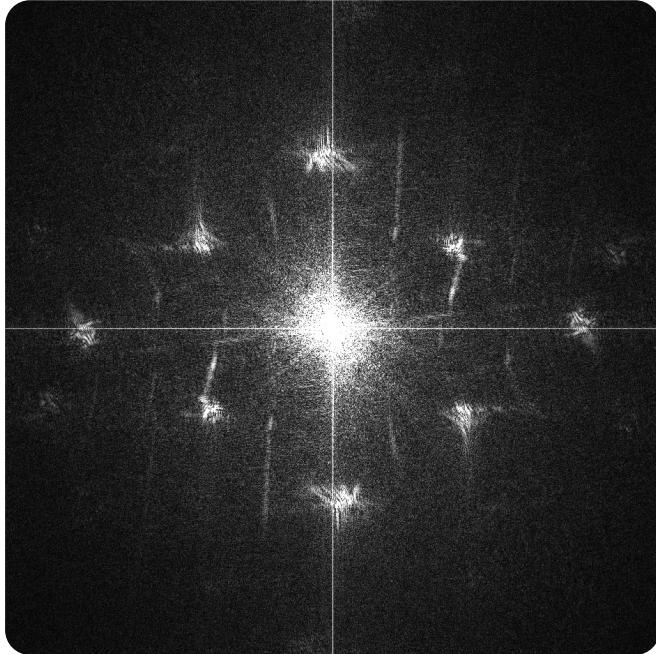


Scaling in Two Dimensions

- Just do one dimension at a time
- Either order produces same result (separable)

Terminology

- Support of f is interval where it is nonzero
 - Can be finite (e.g. box) or infinite (e.g. sinc)
- Images typically have finite support
 - So scaling up isn't usually a problem
 - And sinc filters *preserve* detail
- Bilinear filtering isn't as good at this
 - But textures are more often scaled *down*

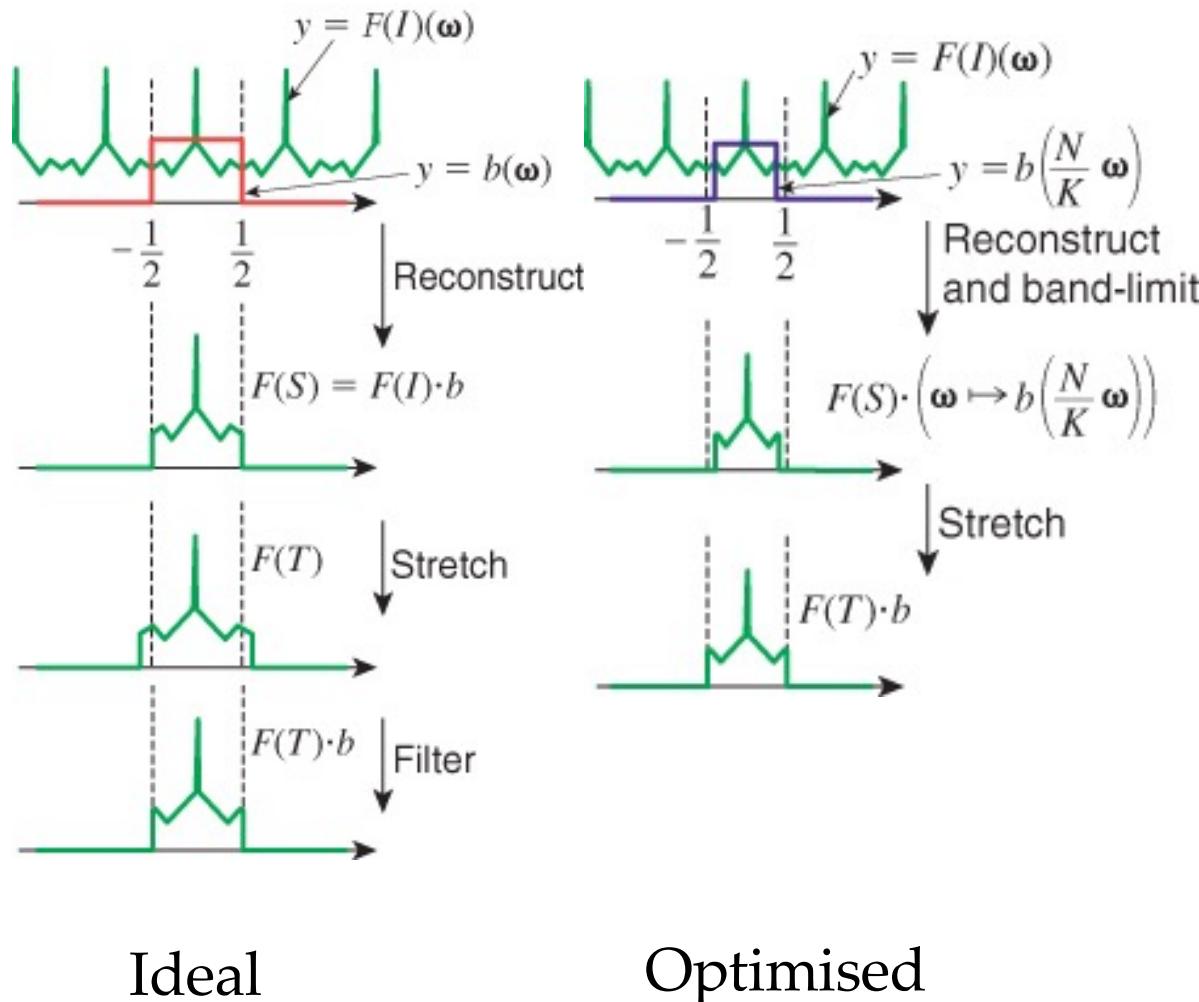


Scale down



- Depends on the method of re-scale
- Limited our band to show higher frequencies
- But can introduce new sharp details

Scaling Down



- We have to add an extra box filter
- In order to preserve the bandlimit



Boundary Effects

- What happens at the image boundary?
- Not enough pixels for filter
- So we have to invent / assume extra pixels
- And each method has it's side effects

Boundary Assumptions

- We can also limit the reconstruction filter
- Or we can cheat & re-weight

Extend by reflection



Original



Extend by zeroes



Side-Effects

- Extend by Zeroes:
 - Leads to darkening around edge of image
- Extend by Reflection:
 - Not necessarily realistic
- Extend by Constants:
 - Tends to smear boundary inwards
 - Finite Support Filter
 - Impossible to do in both domains at once





Band-Limited Filters

- We want a wider filter than linear
- But not as wide as sinc
- Iterate the box (tent) filter multiple times
- Leads to cubic B-spline filters
 - See 5821M for discussion

The Standard Hack

- Each filter assigns weights to pixels
- And *should* guarantee they sum to 1
 - (otherwise we lose/gain brightness overall)
- Omit filter pixels outside the image
 - This drops the sum below 1
 - So divide by the sum of the weights kept
- Result: interpolation only uses valid pixels



Images by:

- S2: Matt moloney
- S3: Kalea Jerielle
- S6: Zachary Keimig
- S40: Nathan dumlao

